




# 14013204-3 - PARALLEL COMPUTING



## Chapter 2

---

## Parallel Hardware


# Roadmap

- 2.1 Some background
  - 2.1.1 The von Neumann architecture
  - 2.1.2 Processes, multitasking, and threads
- 2.2 Modifications to the von Neumann model
  - 2.2.1 The basics of caching
  - 2.2.2 Cache mappings
  - 2.2.3 Caches and programs: an example
  - 2.2.4 Virtual memory
  - 2.2.5 Instruction-level parallelism
  - 2.2.6 Hardware multithreading
- 2.3 Parallel hardware
  - 2.3.1 Classifications of parallel computers
  - 2.3.2 SIMD systems
  - 2.3.3 MIMD systems

## **2.1 Some Background**


### **The von Neumann architecture**

# SOME BACKGROUND



- Parallel hardware and software have grown out of conventional serial hardware and software: hardware and software that runs (more or less) a single job at a time.
- So to better understand the current state of parallel systems, let's begin with a brief look at a few aspects of serial systems.

# Serial hardware and software



# The von Neumann Architecture

- For over 40 years, virtually all computers have followed a **common machine model** known as the **von Neumann computer**. Named after the Hungarian mathematician John von Neumann.
- The “classical” von Neumann architecture consists of **main memory**, a **central processing unit (CPU)** or processor or core, and an **interconnection between the memory and the CPU**.
- A von Neumann computer uses the **stored-program** concept. The CPU executes a stored program that specifies a sequence of **read and write** operations on the memory.
- A von Neumann machine executes a **single instruction at a time**, and each instruction operates on only a few pieces of data.





Figure 2.1

# Main memory

- This is a collection of locations, each of which is capable of **storing both instructions and data**.
- Every location consists of an **address**, which is used to access the location, and the contents of the location.



# Central processing unit (CPU)


- Divided into two parts.
  1. **Control unit** - responsible for deciding which instruction in a program should be executed. (*the boss*)
  2. **Datapath** - responsible for executing the actual instructions. (*the worker*)
    - Combination of **different functional units** including ALU, registers, and internal buses that work together to process the instructions and manage the data.
    - The ALU is a critical component of the datapath responsible for performing **arithmetic and logical operations** on data. It executes calculations based on the instructions provided by the CPU's control unit.




## Key terms

- **Register** – very fast storage, part of the CPU.
  - Store Data in the CPU and information about the state of an executing program.
- **Program counter** – special register in the control unit - stores the address of the next instruction to be executed.
- **Bus** – wires, and hardware that connect the CPU and memory.
  - **Instructions and data** are transferred between the CPU and memory via the interconnect.
  - This has traditionally been a bus, which consists of a collection of parallel wires and some hardware controlling access to the wires. More recent systems use more complex interconnects.

memory




fetch/read



CPU


- A von Neumann machine **executes a single instruction at a time**, and each instruction operates on only a few pieces of data.
- **When data or instructions are transferred from memory to the CPU**, we sometimes say the data or instructions are **fetched** or **read** from memory.



- When data are transferred from the CPU to memory, we sometimes say the data are **written** to memory or **stored**.

# von Neumann bottleneck

- The separation of memory and CPU is often called the **von Neumann bottleneck** since **the interconnect determines the rate at which instructions and data can be accessed.**
- The potentially vast quantity of data and instructions needed to run a program is effectively isolated from the CPU.
- In 2021, CPUs are capable of executing instructions **more than one hundred times faster** than they can fetch items from main memory.



## **2.1 Some Background**

### **Processes, multitasking, and threads**

# An operating system “process”



- Recall that the **operating system (OS)** is a major piece of software, whose purpose is to **manage hardware and software resources** on a computer.
- An OS determines which programs can run and when they can run.
- It also controls the **allocation of memory** to running programs and **access to peripheral devices** such as hard disks and network interface cards.
- When a user runs a program, the operating system creates a **process**.
- A **process** is an instance of a computer program that is being executed.

# An operating system “process”

- Components of a process:
  - The executable machine language program.
  - A block of memory: which will include the **executable code**, a **call stack** that keeps track of active functions, a **heap** that can be used for memory explicitly allocated by the user program, and some other memory locations.
  - Descriptors of resources the OS has allocated to the process e.g. file descriptors.
  - Security information e.g. information specifying which hardware and software resources the process can access.
  - Information about the state of the process - such as whether the process is ready to run or is waiting on some resource, the content of the registers, and information about the process's memory.

# An operating system “process”

- As a process executes, it changes its **state**
  - new**: The process is being created
  - running**: Instructions are being executed
  - waiting**: The process is waiting for some event to occur
  - ready**: The process is waiting to be assigned to a CPU
  - terminated**: The process has finished execution



# Multitasking

- Most modern operating systems are **multitasking**.
- Gives the illusion that a single processor system is running multiple programs simultaneously.
- Each process takes turns running. (**time slice**)
- After its time is up, it waits until it has a turn again (**blocks**) and the operating system can run a different program.
- A multitasking OS may **change the running process many times in a minute**, even though changing the running process can take a long time.
- In a multitasking OS, if a process needs to wait for a resource—for example, it needs to read data from external storage—it will **block**. This means that it will stop executing and the operating system can run another process.
  - However, many programs can continue to do useful work even though the part of the program that is currently executing must wait on a resource. For example, an airline reservation system that is blocked waiting for a seat map for one user could provide a list of available flights to another user.


# Threading

- **Threading** provides a mechanism for programmers to **divide their programs into more or less independent tasks** with the property that when one thread is blocked another thread can be run.
- Furthermore, in most systems it's possible to switch between threads much **faster** than it's possible to switch between processes.
- This is because threads are “**lighter weight**” than processes.
- **Threads** are contained within processes, so they can **use the same executable**, and they usually **share the same memory and the same I/O devices**.

# Threading


- In fact, two threads belonging to one process can share most of the process' resources.
- The two most important exceptions are:
  1. They'll need a record of **their own program counters**
  2. And they'll need **their own call stacks** so that they can execute independently of each other.

# A process and two threads




## 2.2 Modifications To The Von Neumann Model

- As we noted earlier since the first electronic digital computers were developed back in the 1940s, computer scientists and computer engineers have made **many improvements to the basic von Neumann architecture.**
- Many are targeted at **reducing the problem of the von Neumann bottleneck**, but many are also targeted at simply **making CPUs faster**.
- We will look at three of these improvements:
  1. **caching**
  2. **virtual memory**
  3. **low-level parallelism.**



# Modifications To The Von Neumann Model

## 1. caching



# Basics of caching

- A collection of memory locations that can be accessed in less time than some other memory locations.
- A **CPU cache** is a collection of memory locations that the CPU can access **more quickly** than it can access main memory.
- A CPU cache can either be located on the same chip as the CPU or it can be located on a separate chip that can be accessed much faster than an ordinary memory chip.



# Basics of Caching – Related Issues

1. **Block Placement:** Where to Place the Main Memory Block in the Cache?
  2. **Block Identification:** How to Find the Main Memory Block in the Cache?
  3. **Block Replacement:** During the Cache Miss, how to choose which entry to replace from the cache?
  4. **Write Strategy:** How are the updates propagated? – *will be covered in the next lecture with other issues related to cache coherency and false sharing*
- You have already studied the first three points before. Watch these videos to review:  
<https://www.youtube.com/watch?v=1tvW8kzOpSA&list=PLBlnK6fEyqRgLLIzdgiTUKULKJPYc0A4q&index=25>
  - <https://www.youtube.com/watch?v=7IxAfszjy68&list=PLBlnK6fEyqRgLLIzdgiTUKULKJPYc0A4q&index=26>
  - [https://www.youtube.com/watch?v=\\_Hh-NcdbHCY&list=PLBlnK6fEyqRgLLIzdgiTUKULKJPYc0A4q&index=27](https://www.youtube.com/watch?v=_Hh-NcdbHCY&list=PLBlnK6fEyqRgLLIzdgiTUKULKJPYc0A4q&index=27)
  - <https://www.youtube.com/watch?v=QyQf9KvkQA4&list=PLBlnK6fEyqRgLLIzdgiTUKULKJPYc0A4q&index=28>

# Caches and programs

```
double A[MAX][MAX], x[MAX], y[MAX];
```

```
/* Initialize A and x, assign y = 0 */
```

```
/* First pair of loops */
```

```
for (i = 0; i < MAX; i++)
    for (j = 0; j < MAX; j++)
        y[i] += A[i][j]*x[j];
```

```
/* Assign y = 0 */
```


```
/* Second pair of loops */
for (j = 0; j < MAX; j++)
    for (i = 0; i < MAX; i++)
        y[i] += A[i][j]*x[j];
```

A total of 4 misses

| Cache Line | Elements of A |         |         |         |
|------------|---------------|---------|---------|---------|
| 0          | A[0][0]       | A[0][1] | A[0][2] | A[0][3] |
| 1          | A[1][0]       | A[1][1] | A[1][2] | A[1][3] |
| 2          | A[2][0]       | A[2][1] | A[2][2] | A[2][3] |
| 3          | A[3][0]       | A[3][1] | A[3][2] | A[3][3] |

A total of 16 misses


# **Modifications To The Von Neumann Model**



## **2. virtual memory**


# Virtual memory

- If we run a very large program or a program that accesses very large data sets, all of the instructions and data **may not fit into main memory**.
- **Virtual memory** functions as a cache for secondary storage.



Teach-ICT

# Virtual memory




- It exploits the principle of **spatial and temporal locality**.
- Keeping in **main memory** only the **active parts** of the many running programs; those parts that are idle are kept in a block of secondary storage called **swap space**.
- Like **CPU caches**, virtual memory operates on **blocks of data and instructions**. These blocks are commonly called **pages**, and since secondary storage access can be hundreds of thousands of times slower than main memory access, pages are relatively large (page size ranges from 4 to 16 KB).

# Virtual memory

- The relative slowness of disk accesses has a couple of additional consequences for virtual memory.
  - We as programmers don't directly control virtual memory.
  - Unlike CPU caches, which are handled by system hardware, virtual memory is usually controlled by a combination of system hardware and operating system software.
- You should have studied the virtual memory management details in OS course. To review: Read the book (pages: 25- 27)

# Modifications To The Von Neumann Model



## 3. low-level parallelism.

# Instruction Level Parallelism (ILP) (1)

- **Instruction-level parallelism (ILP)** attempts to improve processor performance by having multiple processor components or **functional units** simultaneously executing instructions in parallel within a single program or thread.
- It **exploits parallelism at the instruction level**.
- There are two approaches to instruction-level parallelism:
  1. **Hardware**: hardware level works upon **dynamic parallelism**.
    - Dynamic parallelism means the **processor** decides at run time which instructions to execute in parallel.
  2. **Software**: software level works on **static parallelism**.
    - Static parallelism means the **compiler** decides which instructions to execute in parallel.


## Instruction Level Parallelism (ILP) (2)

- There are two main approaches to ILP on hardware level :
  1. **Pipelining** - functional units are arranged in stages.
  2. **Multiple issue** - multiple instructions can be simultaneously initiated.
- Both approaches are used in virtually all modern CPUs.

# Pipelining

The principle of pipelining is similar to a **factory assembly line**:

1. while one team is bolting a car's engine to the chassis.
2. another team can connect the transmission to the engine and the driveshaft of a car that's already been processed by the first team.
3. and a third team can bolt the body to the chassis in a car that's been processed by the first two teams.



# Pipelining example (1)

- suppose we want to add the floating-point numbers  $9.87 \times 10^4$  and  $6.54 \times 10^3$ .  
Then we can use the following steps:

| Time | Operation         | Operand 1          | Operand 2           | Result               |
|------|-------------------|--------------------|---------------------|----------------------|
| 1    | Fetch operands    | $9.87 \times 10^4$ | $6.54 \times 10^3$  |                      |
| 2    | Compare exponents | $9.87 \times 10^4$ | $6.54 \times 10^3$  |                      |
| 3    | Shift one operand | $9.87 \times 10^4$ | $0.654 \times 10^4$ |                      |
| 4    | Add               | $9.87 \times 10^4$ | $0.654 \times 10^4$ | $10.524 \times 10^4$ |
| 5    | Normalize result  | $9.87 \times 10^4$ | $0.654 \times 10^4$ | $1.0524 \times 10^5$ |
| 6    | Round result      | $9.87 \times 10^4$ | $0.654 \times 10^4$ | $1.05 \times 10^5$   |
| 7    | Store result      | $9.87 \times 10^4$ | $0.654 \times 10^4$ | $1.05 \times 10^5$   |

- Here we're using base 10 and a three-digit mantissa or significand with one digit to the left of the decimal point.
- Normalizing** shifts the decimal point one unit to the left, and **rounding** rounds to three digits.
- If each of the operations takes **one nanosecond** ( $10^{-9}$  seconds), the addition operation will take **seven nanoseconds**.

## Pipelining example (2)

```
float x[1000], y[1000], z[1000];  
.  
.  
.  
for (i = 0; i < 1000; i++)  
    z[i] = x[i] + y[i];
```

- This for loop takes about **7000 nanoseconds**.

## Pipelining example (3)

- As an alternative, suppose we **Divide** the floating-point adder into 7 separate pieces of **hardware or functional units**.
  - First unit fetches two operands, the second unit compares exponents, etc.
  - Output of one functional unit is input to the next.
- Then a single floating-point addition will still take 7 nanoseconds.
- However, when we execute the for loop, we can **fetch  $x[1]$**  and  **$y[1]$**  while we're **comparing the exponents of  $x[0]$  and  $y[0]$** .

## Pipelining example (4)

- More generally, it's possible for us to **simultaneously execute seven different stages in seven different additions**


| Time | Fetch | Compare | Shift | Add | Normalize | Round | Store |
|------|-------|---------|-------|-----|-----------|-------|-------|
| 0    | 0     |         |       |     |           |       |       |
| 1    | 1     | 0       |       |     |           |       |       |
| 2    | 2     | 1       | 0     |     |           |       |       |
| 3    | 3     | 2       | 1     | 0   |           |       |       |
| 4    | 4     | 3       | 2     | 1   | 0         |       |       |
| 5    | 5     | 4       | 3     | 2   | 1         | 0     |       |
| 6    | 6     | 5       | 4     | 3   | 2         | 1     | 0     |
| :    | :     | :       | :     | :   | :         | :     | :     |
| 999  | 999   | 998     | 997   | 996 | 995       | 994   | 993   |
| 1000 |       | 999     | 998   | 997 | 996       | 995   | 994   |
| 1001 |       |         | 999   | 998 | 997       | 996   | 995   |
| 1002 |       |         |       | 999 | 998       | 997   | 996   |
| 1003 |       |         |       |     | 999       | 998   | 997   |
| 1004 |       |         |       |     |           | 999   | 998   |
| 1005 |       |         |       |     |           |       | 999   |

Table 2.3: Pipelined Addition.

Numbers in the table are subscripts of operands/results.

## Pipelining example (5)

- One floating point addition still takes **7 nanoseconds**, But 1000 floating point additions now takes **1006 nanoseconds!**
- In general, **a pipeline with k stages won't get a k-fold improvement in performance.**
  - For example, if the times required by the various functional units are different, then the stages will effectively run at the speed of the slowest functional unit.
  - Furthermore, delays—such as waiting for an operand to become available—can cause the pipeline to stall.




# Multiple Issue (1)

- Pipelines improve performance by **taking individual pieces of hardware or functional units and connecting them in sequence**.
- Multiple issue processors **replicate functional units** and try to **simultaneously execute different instructions** in a program.
  - if we have **two complete floating point adders**, we can approximately **halve the time** it takes to execute the loop.
  - While the first adder is computing  $z[0]$ , the second can compute  $z[1]$ ; while the first is computing  $z[2]$ , the second can compute  $z[3]$ ; and so on.

**for (i = 0; i < 1000; i++)**

$$z[i] = x[i] + y[i];$$



# Multiple Issue (2)

- If the functional units are scheduled at compile time, the multiple issue system is said to use **static multiple issue**.
- If they're scheduled at run-time, the system is said to use **dynamic multiple issue**.
- A processor that supports dynamic multiple issue is sometimes said to be **superscalar**.
- Superscalar architecture is a method of designing processors to **allow multiple instructions to be issued and executed by multiple execution units in a single clock cycle**.
- It's based on the concept of instruction-level parallelism (ILP).
- In a superscalar processor, there are **multiple execution units** such as **arithmetic logic units (ALUs)**, **floating-point units (FPUs)**, and more. These units can work concurrently, executing different instructions from a single instruction stream.

# Speculation (1)

- In order to make use of multiple issue, the system must find instructions that can be executed simultaneously.
- One of the most important techniques is **speculation**.



- In speculation, the compiler or the processor makes a guess about an instruction and then executes the instruction on the basis of the guess.


## Speculation (2)

```
z = x + y;  
if ( z > 0 )  
    w = x;  
else  
    w = y;
```



If the system speculates incorrectly,  
it must go back and recalculate  $w = y$ .

```
z = x + y;  
w = *a_p;
```



If the system speculates incorrectly,  
it must go back and re-execute  $w = *a_p$ .

## Speculation (3)

- If the **compiler** does the speculation, it will usually **insert code that tests whether the speculation was correct**, and, if not, takes corrective action.
- If the **hardware** does the speculation, the processor usually **stores the result(s) of the speculative execution in a buffer**.
  - When it's known that the **speculation was correct**, the **contents of the buffer are transferred to registers or memory**.
  - If the speculation was **incorrect**, the **contents of the buffer are discarded**, and the **instruction is re-executed**.

## Thread Level Parallelism (TLP)

- ILP can be very difficult to exploit: a program with a **long sequence of dependent statements** offers few opportunities.
- For example, in a direct calculation of the Fibonacci numbers

$f[0] = f[1] = 1;$

for ( $i = 2; i \leq n; i++$ )

$f[i] = f[i-1] + f[i-2];$

- There's essentially **no opportunity** for simultaneous execution of instructions.

# Thread Level Parallelism (TLP)

- **Thread-level parallelism (TLP)** : attempts to provide parallelism through the simultaneous execution of **different threads**, so it provides **coarser-grained parallelism** than **ILP**,
- That is, the program units that are being simultaneously executed—threads—are **larger or coarser** than the **finer-grained** units—individual instructions.

# Hardware multithreading (1)

- **Hardware multithreading** provides a means for systems to continue doing useful work when the task being currently executed has stalled.
  - Ex., the current task has to wait for data to be loaded from memory
- Instead of looking for parallelism in the currently executing thread, it may make sense to simply run another thread.
- Of course, for this to be useful, the system must support very rapid switching between threads.

## Hardware multithreading (2)

- **Fine-grained multithreading** - the processor switches between threads after each instruction, skipping threads that are stalled (*experienced a delay or pause in execution* ).
  - Pros: potential to avoid wasted machine time due to stalls.
  - Cons: a thread that's ready to execute a long sequence of instructions may have to wait to execute every instruction.

# Hardware multithreading (3)


- **Coarse-grained multithreading** - only switches threads that are stalled waiting for a time-consuming operation to complete (e.g., a load from main memory).
    - Pros: Switching threads doesn't need to be nearly instantaneous.
    - Cons: The processor can be idled on shorter stalls, and thread switching will also cause delays.

## Hardware multithreading (4)

- **Simultaneous multithreading (SMT)** - a variation on fine-grained multithreading.
- It attempts to exploit superscalar processors by allowing multiple threads to make use of the multiple functional units.

# Hardware visible to programmer

- Multiple issue, Pipelining can clearly be considered to be **parallel hardware**, since functional units can be executed simultaneously
- This form of **parallelism** is **not** usually **visible** to the programmer.
  - We're treating them as extensions to the basic **von Neumann model**
- we will talk about **parallel hardware** that is **visible** to the programmer (must modify the source code to exploit it) → **Flynn's taxonomy** .




A programmer can write code to exploit.

## 2.3 Parallel hardware

# Basic Parallel Architecture

- Still consist mainly of the same three main building blocks:
  - **The memory modules:** to store programs and their data.
  - **The interconnection network:** to communicate with other processors, and to communicate with memory.
    - Fetch/Read
    - Write/Store
  - **The processors (the cores): more than one core** to execute the parallel program.

# Basic Parallel Architecture



# Basic Parallel Architecture

- Over recent decades, there has been a gradual development in the degree of sophistication of each of these building blocks.
- What makes one parallel computer different from another is **the manner in which they are arranged.**
- Parallel computing cores themselves are essentially similar to those used in single-core systems with new technologies increasingly embedded with the cores for parallel architectures
  - The aim is to add more improvements or new functionalities e.g., GPUs, and FPGAs.

# Classifications of parallel computers

- **Parallel computers** have been classified in different ways.
- The earliest and the most widely used classification is Flynn's Taxonomy, 1967 ([Flynn, Michael. \(1967\). Very High-Speed Computing Systems. Proceedings of the IEEE. 54. 1901 - 1909. 10.1109/PROC.1966.5273.](#))
- Flynn's taxonomy distinguishes multi-processor computer architectures according to how they can be classified along the two independent dimensions of **Instruction** and **Data**.
- Each of these dimensions can have only one of two possible states: **Single** or **Multiple**.



# Flynn's Taxonomy

- Flynn classified computing architectures into four categories, based on the number of **instruction streams** and **data streams** they can simultaneously manage:

|  |  |
|--|--|
| <i>classic von Neumann</i><br><b>(SISD)</b><br><b>Single instruction stream</b><br><b>Single data stream</b> | <b>(SIMD)</b><br><b>Single instruction stream</b><br><b>Multiple data stream</b>   |
| <b>(MISD)</b><br><b>Multiple instruction stream</b><br><b>Single data stream</b><br><i>not covered</i>       | <b>(MIMD)</b><br><b>Multiple instruction stream</b><br><b>Multiple data stream</b> |



# SISD

- It is a **sequential** computer (Not parallel) with a single processor → **classical von Neumann system**.
- **Single instruction**: Only one instruction stream is being acted on by the CPU during any **one clock cycle**.
- **Single data**: Only one data stream is being used as **input** during any one clock cycle.
- Deterministic execution.
- This is the oldest and until recently, the most **prevalent** form of computer.




# SIMD

- Parallelism is achieved by dividing data among the processors.
- Applies the same **single instruction** to **multiple data items**.
- An abstract SIMD system can be thought of as having **a single control unit and multiple datapaths**.
  - An instruction is broadcast **from the control unit to the datapaths**, and each datapath either applies the instruction to the current data item, or it is idle.
  - Support **data parallelism**.



# SIMD example



```
for (i = 0; i < n; i++)
```

```
    x[i] += y[i];
```

# SIMD

- What if we don't have as many Datapaths as data items?
- Divide the work and process iteratively.
- Ex.  $m = 4$  Datapaths and  $n = 15$  data items.

| Round | Datapath <sub>1</sub> | Datapath <sub>2</sub> | Datapath <sub>3</sub> | Datapath <sub>4</sub> |
|-------|-----------------------|-----------------------|-----------------------|-----------------------|
| 1     | X[0]                  | X[1]                  | X[2]                  | X[3]                  |
| 2     | X[4]                  | X[5]                  | X[6]                  | X[7]                  |
| 3     | X[8]                  | X[9]                  | X[10]                 | X[11]                 |
| 4     | X[12]                 | X[13]                 | X[14]                 |                       |

# SIMD drawbacks

- All Datapaths are required to execute the same instruction or remain idle → which can seriously **degrade the overall performance**.
  - For example, suppose we only want to carry out the addition if  $y[i]$  is positive:

```
for ( i = 0; i < n ; i ++)
    if ( y [ i ] > 0 . 0 ) x [ i ] += y [ i ] ;
```
  - Some Datapaths will be idle if the condition is not true while others can proceed with the computation.
- In classic design, they must also operate **synchronously**, that is, each datapath **must wait** for the next instruction to be broadcast before proceeding.
- The Datapaths **have no instruction storage**, so a datapath can't delay the execution of an instruction by storing it for later execution.
- Efficient for large **data parallel** problems, but not other types of more complex parallel problems.

# SIMD Systems


- SIMD systems have had a somewhat checkered history.
  - In the early 1990s, a maker of SIMD systems (Thinking Machines) was one of the largest manufacturers of parallel supercomputers.
  - However, by the late 1990s, the only widely produced SIMD systems were **vector processors**.
  - More recently, **graphics processing units (GPUs)**, and desktop CPUs are making use of aspects of SIMD computing.

# Vector processors (1)

- Operate on **arrays or vectors** of data while conventional CPUs operate on **individual data elements or scalars**.
  - The operand to the instructions are complete vectors instead of one element
- Typical recent systems have the following characteristics:
  1. **Vector registers:** Capable of storing a **vector of operands** and operating simultaneously on their contents.
    - The vector length is fixed by the system and can range from 4 to 256 64-bit elements.
  2. **Vectorized and pipelined functional units:** The **same operation** is applied to each element in the vector, or, in the case of operations like addition, the same operation is applied to each pair of corresponding elements in the two vectors. → **vector operations are SIMD.**


# Vector processors (2)

**SCALAR**  
**(1 operation)**



**add r3, r1, r2**

**VECTOR**  
**(N operations)**



**add.vv v3, v1, v2**

# Vector processors (3)

## 3. Vector instructions.

- Operate on vectors rather than scalars.
- If the vector length is `vector_length`, these instructions have the great virtue that a simple loop such as

```
for ( i = 0; i < n ; i ++ )
```

```
    x [ i ] += y [ i ] ;
```

requires only a single load, add, and store for each block of `vector_length` elements, while a conventional system requires a load, add, and store for each element.

# Vector processors (3)

## 4. optimized memory access

- Interleaved memory.
  - Multiple “banks” of memory, which can be accessed more or less independently.
  - Distribute elements of a vector across multiple banks, so reduce or eliminate delay in loading/storing successive elements.
- Strided memory access
  - The program accesses elements of a vector located at fixed intervals.

# Vector processors - Pros

- Fast.
- Easy to use.
- Vectorizing compilers are good at identifying code to exploit.
- Compilers also can provide information about code that cannot be vectorized.
  - Helps the programmer re-evaluate code.
- High memory bandwidth.
- Uses every item in a cache line.




## Vector processors - Cons

- They don't handle irregular data structures as well as other parallel architectures.
- A very finite limit to their ability to handle ever larger problems. (scalability)



# Graphics Processing Units (GPU)

- Real time graphics application programming interfaces or API's use points, lines, and triangles to internally represent the surface of an object.



# GPUs

- A **graphics processing pipeline** converts the internal representation into an array of pixels that can be sent to a computer screen.
- Several stages of this pipeline (called **shader functions**) are programmable.
  - Typically just a few lines of C code.




# GPUs

- Shader functions are also implicitly parallel, since they can be applied to multiple elements in the graphics stream.
- GPU's can often optimize performance by using SIMD parallelism.
- The current generation of GPUs use SIMD parallelism.
  - Although they are not pure SIMD systems.
- GPUs are becoming increasingly popular for general, high-performance computing, and several languages have been developed that allow users to exploit their power.

# MISD



- It is not very common. Mainly mentioned for completeness.
- Multiple functional units or processors execute **multiple instructions** concurrently on the same **single data stream**.
  - A **single data stream** is **fed into multiple processing units**.
  - Each processing unit operates on the data independently via independent instruction streams.



Source: <https://en.wikipedia.org/wiki/MISD>

# MIMD

- Currently, the **most common type of parallel computer**. Most modern computers fall into this category.
- **Multiple Instruction**: Every processor may be executing a different instruction stream
- **Multiple Data**: Every processor may be working with a different data stream



Source: <https://en.wikipedia.org/wiki/MIMD>


# MIMD

- **MIMD systems** typically consist of a collection of fully independent processing units or cores, each of which has its own control unit and its own Datapath.
- Furthermore, unlike SIMD systems, MIMD systems are usually **asynchronous**, that is, the processors can operate at their own pace.
- Factors that promote the development of MIMD architectures:
  - Computers in this category offer support for a wider range of parallel algorithms  
→ with wider applicability.
  - MIMD is extremely cost-effective
- We will focus our study on MIMD architectures ..


# Classification of MIMD memory Architectures

- Two main categories of MIMD architectures, based on the organization of memory:
  1. **Shared memory architectures** → all cores **share access** to the **same** memory.
  2. **Distributed memory architectures** → each core has its **own** local memory, and accesses values in other memories using the network.
    - the processors usually communicate explicitly by **sending messages** or by using special functions that provide access to the memory of another processor.

# Classification of MIMD memory Architectures



**Shared memory architectures**



**Distributed memory architectures**